



Mathematical Fragility: An Investigation on Birth and Death Chains

Bailey Bjornstad



August 29, 2017



Contents

1	Project Update - Mid August	1
1.1	Birth and Death Chains: Recurrence	1
1.2	Chains with Variable Step Size, and an Investigation of Hitting Times	8

Chapter 1

Project Update - Mid August

1.1 Birth and Death Chains: Recurrence

As a subset of Markov chains, birth and death chains are stochastic processes in which the probability distribution is dependent on the current position. The following diagram enlightens this, insert diagram here. This can be thought of as a fog obscuring a point of interest (say, the origin) with density given by $p(x)$. It is reasonable to consider the "fragility" of this system through the lens of recurrence, a mathematical property of Markov chains that categorizes states of a chain based on whether the chain is expected to return to that state. Further, in the case of chains with infinite states—a birth and death chain where the states are values of \mathbb{Z} , for instance—a distinction is drawn between those states where the expected return time is infinite, and those where the expected return time is finite. States are thus called *null-recurrent* or *positive-recurrent*, respectively. In contrast, a state is called *transient* if the chain cannot be expected to return.

This allows a mathematical characterization of instances in which the chain seems to "get lost" from the origin. For instance, if the origin is classified as a transient state, then a walker starting at the origin cannot be expected to return to the origin (he may, but not with full probability). In contrast, if the origin were classified as a recurrent state, then the walker will return with probability 1. Further, if the origin were null-recurrent, the walker may find himself diverging from the origin for arbitrary periods of time, but will with probability 1 eventually return. If the origin were positive-recurrent, then a concrete estimate can be placed on the walker's return time to the origin. This is to say that transient states are the "most lost," as there can be no concrete expectations about a walker's return to such states, nor can concrete estimates be made about how long returning may take. Null-recurrent states are still "somewhat lost," though less so than transient states due to the fact that concrete expectations *can* be made about a walker's return. Positive-recurrent states are the "least lost" since a walker can be expected to return within a finite amount of time.

In order to investigate how recurrence properties impact a chain's "fragility," I first needed to analyze some specific examples of birth and death chains. In particular, I defined a chain in which the state-space—the possible positions that the chain can take—is the integers. Then,

let

$$p(x) = \begin{cases} \frac{1}{2+|x|^{-\beta}} & \text{if } x > 0 \\ 1 - \frac{1}{2+|x|^{-\beta}} & \text{if } x < 0 \\ \frac{1}{2} & \text{if } x = 0 \end{cases}$$

for some $\beta \geq 0$, and let $q(x) = 1 - p(x)$. Let $r(x) = 0$. This represents a birth and death chain that is somewhat biased towards the origin, and in which the limiting probabilities

$$\lim_{x \rightarrow \infty} p(x) = \lim_{x \rightarrow \infty} q(x) = \frac{1}{2}$$

and


$$\lim_{x \rightarrow -\infty} p(x) = \lim_{x \rightarrow -\infty} q(x) = \frac{1}{2}.$$

Physically, this should represent that a walker at extremely far distances from the origin cannot see it, and so should step either in the positive direction or in the negative direction with probability $1/2$. We will refer to this birth and death chain as the *power-distributive birth and death chain* with parameter β .


Unfortunately, unlike in the case of stationary Markov chains in which probabilities are independent of the position—which usually simplifies analysis—the recurrence properties of this birth and death chain will need to be investigated a different way. I hypothesize, in particular, that if $\beta < 1$, then the origin in the above chain will be a positive-recurrent state, and that if $\beta > 1$, then the origin will be a null-recurrent state. In order to develop a stronger intuition about this behavior, I turned to some numerical simulations in Python, using the NumPy and Matplotlib libraries.

I began by writing a base-class that represents a random walk, from which more specific instances can be derived. I then developed a birth and death class inherited from the random walk base class. This implemented the specific probability distribution detailed above. Then, for each β in a range of β values—defined by a minimum β , maximum β , and β step-size—a simulation of one-million time steps was run, with the origin as the initial position, and the given probability distributions. Then, using the fact that storage of return times is implemented in the random walk class, the mean and standard-deviation of the return times was calculated, and plotted using Matplotlib's pyplot class. In particular, I used the scatter-plot, histogram, and 2-dimensional histogram functions.

For example, the following plot shows the mean return times of a simulation with β values ranging from 0 to 3 with a step size of 0.005.



Here, it appears as though the mean return times begin to scatter into higher values shortly after β grows beyond 1, thus supporting that the chains may be getting "more lost" more frequently. This suggests that the origin in the given birth and death chain may be either null-recurrent or transient when $\beta > 1$ and is, in contrast, positive-recurrent when $\beta < 1$. The corresponding 2-dimensional histogram similarly supports this conclusion.



However, this is not evidence enough to support the claim, especially in mathematics where the standard of proof lies beyond statistical inference. As such, I turned my attention toward

proving the conjecture, using a variety of analytical tools. According to Kobayashi's *Probability, Random Processes, and Statistical Analysis*, we mathematically define a recurrent or transient state the following way. Begin by defining the *first-passage time* T_{ij} as the number of time steps it takes for the chain to transition from state i to state j . Then, we let $\rho_{ij}^{(n)}$ be the probability that $T_{ij} = n$; in other words, $\rho_{ij}^{(n)} = P(T_{ij} = n)$. Thus, the following sum

$$\rho_{ij} = \sum_{n=1}^{\infty} \rho_{ij}^{(n)}$$

represents the probability that the chain ever starts at state i and transitions to state j . Of particular importance is the quantity ρ_{ii} , the probability that the chain ever returns to state i after first beginning there. From this, we can classify states.

Definition 1.1.1 A state i in the state space \mathcal{S} is *transient* if $\rho_{ii} < 1$ and is *recurrent* if $\rho_{ii} = 1$.

However, computation of ρ_{ii} for any fixed state i proved difficult with the distribution given above. Unfortunately, the position-dependence of $p(x)$ and $q(x)$ make combinatoric approaches highly challenging, and I had to turn my attention elsewhere.

Observe, first, that $p(x) = q(-x)$ when $x > 0$. As a result, the birth and death chain is symmetric about the origin. Further, since states change in integer increments, if the chain begins to observe positive values, it must do so at least until the chain returns to the origin. Due to these two facts, the given birth and death chain can be considered as if it had state space \mathbb{N} instead of \mathbb{Z} . Thus, it suffices to consider the birth and death chain on \mathbb{N} with probability distribution given by

$$p(x) = \begin{cases} 1 & \text{if } x = 0 \\ \frac{1}{2+x^{-\beta}} & \text{if } x > 0 \end{cases},$$

and $q(x) = 1 - p(x)$. In particular, $p(0) = 1$ since $r(0) = 0$ and $q(0) = 0$ by definition.

Importantly, a new concept must be introduced: that of *communication* between states. The definition comes from Kobayashi.

Definition 1.1.2 Let i and j be states. State j is *reachable* from state i if there exists some $m \in \mathbb{N}$ so that the quantity $P_{ij}^{(m)} = \mathbb{P}[X_{n+m} = j \mid X_n = i] > 0$, and is denoted $i \rightarrow j$. If $i \rightarrow j$ and $j \rightarrow i$, then the states i and j *communicate*, denoted $i \leftrightarrow j$.

It is a common fact of Markov chains that the relation defined by communication is an equivalence relation on the state space. Further, it is another common fact that all states in a particular equivalence class defined by communication must be of the same type. This is to say that if i and j belong to the same equivalence class \mathcal{C} , and i is a recurrent state, then j is also a recurrent state (or vice-versa). Further, we define the concept of *irreducibility*.

Definition 1.1.3 Let \mathcal{C} be a set of states. \mathcal{C} is *irreducible* if $i \leftrightarrow j$ for every i and j in \mathcal{C} .

I claim that the birth and death chain with the given probability distribution is an irreducible chain.

Proposition 1.1.1 The birth and death chain with probability distribution given by

$$p(x) = \begin{cases} 1 & \text{if } x = 0 \\ \frac{1}{2+x^{-\beta}} & \text{if } x > 0 \end{cases},$$

$q(x) = 1 - x$ and $r(x) = 0$ is irreducible.

Proof Note that for every state i with $i > 0$, the quantities $p(i)$ and $q(i)$ are both positive. Now, fix arbitrary states i and j , and without loss of generality suppose that $j > i$ and that $i \neq 0$ and $j \neq 0$. We wish to show that there exists some m so that $P_{ij}^{(m)} = \mathbb{P}[X_{n+m} = j \mid X_n = i]$. Let $m = j - i$. Then,

$$P_{ij}^{(m)} = p(i)p(i+1)p(i+2) \cdots p(i+(j-i)).$$

However, since all terms in this product are positive, so must be the product, so that $P_{ij}^{(m)} > 0$, so that $i \rightarrow j$. Similarly,

$$P_{ji}^{(m)} = q(j)q(j-1)q(j-2) \cdots q(j-(j-i)),$$

which is also positive, so that $j \rightarrow i$. Thus, $i \leftrightarrow j$ when i and j are positive. Now, suppose that $i = 0$. Since it has been determined that $1 \leftrightarrow j$, it suffices to show that $0 \leftrightarrow 1$. Trivially, $p(0) = 1$ and $q(1) = 2/3$, so that $0 \leftrightarrow 1$. Since \leftrightarrow defines an equivalence relation, it follows that $0 \leftrightarrow j$ for any j . Thus, the chain must be irreducible.

Since the chain is irreducible, every state in the chain must share the same recurrence-classification. This is to say that all states are either transient, null-recurrent, or positive-recurrent. This allows the use of the following fact, from Zhong Li.

Proposition 1.1.2 An arbitrary birth and death chain on \mathbb{N} with probability distribution defined by $p(x)$ and $q(x)$ is transient if and only if

$$\sum_{x=1}^{\infty} \frac{q(1)q(2)q(3) \cdots q(x)}{p(1)p(2)p(3) \cdots p(x)} < \infty.$$

Proof Proof forthcoming.

Thus, we can prove the following proposition.

Proposition 1.1.3 The power-distributive birth and death chain with parameter β is recurrent for any β .

Proof It suffices to show that the sum

$$\sum_{x=1}^{\infty} \frac{q(1)q(2)q(3) \cdots q(x)}{p(1)p(2)p(3) \cdots p(x)}$$

diverges. Note that for any state i , $q(i) = 1 - p(i)$. Thus,

$$q(i) = 1 - \frac{1}{2 + i^{-\beta}} = \frac{1 + i^{-\beta}}{2 + i^{-\beta}}.$$

Thus, the sum simplifies, so that

$$\sum_{x=1}^{\infty} \frac{q(1)q(2)q(3) \cdots q(x)}{p(1)p(2)p(3) \cdots p(x)} = \sum_{x=1}^{\infty} (1 + 1^{-\beta})(1 + 2^{-\beta}) \cdots (1 + x^{-\beta}).$$

The summand is clearly larger than 1, so that the series must diverge. Thus, the chain is recurrent.

However, we must still show that the chain is positive-recurrent for values of $\beta < 1$ and null-recurrent for values of $\beta \geq 1$. Doing so will involve the use of logarithms and approximations thereof, as well as a key observation about Riemann sums in relation to corresponding integrals. Further, we draw upon a second of Zhong Li's results:

Proposition 1.1.4 An arbitrary birth and death chain on \mathbb{N} with probability distribution defined by $p(x)$ and $q(x)$ is positive-recurrent if and only if

$$\sum_{x=1}^{\infty} \frac{p(0)p(1)p(2) \cdots p(x-1)}{q(1)q(2)q(3) \cdots q(x)} < \infty.$$

Proof Proof Forthcoming

Proposition 1.1.5 The power-distributive birth and death chain with parameter β is positive-recurrent if $\beta < 1$.

Proof Let B denote the value

$$\sum_{x=1}^{\infty} \frac{p(0)p(1)p(2) \cdots p(x-1)}{q(1)q(2)q(3) \cdots q(x)}.$$

It suffices to show that $B < \infty$. Note that,

$$q(x) = 1 - p(x) = \frac{1 + x^{-\beta}}{2 + x^{-\beta}}.$$

Therefore, we can rewrite B as

$$B = \sum_{x=1}^{\infty} \frac{2 + x^{-\beta}}{(1 + 1^{-\beta})(1 + 2^{-\beta}) \cdots (1 + x^{-\beta})}.$$

Now, we may leverage a useful property of logarithms: that they can "convert" products to sums. In particular, it is true that

$$\log((1 + 1^{-\beta})(1 + 2^{-\beta}) \cdots (1 + x^{-\beta})) = \sum_{i=1}^x \log(1 + i^{-\beta}).$$

Further, it is true that for any choice of β and for every $i \geq 1$,

$$\frac{1}{4}i^{-\beta} \leq \log(1 + i^{-\beta}).$$

Thus, we have that

$$\log((1 + 1^{-\beta})(1 + 2^{-\beta}) \cdots (1 + x^{-\beta})) \geq \frac{1}{4} \sum_{i=1}^x i^{-\beta}.$$

Let

$$s(x) = \sum_{i=1}^x i^{-\beta}.$$

We now make the key observation that this sum is a Riemann-sum with rectangular width of 1 approximating the integral of the function $f(i) = i^{-\beta}$. This is to say that

$$s(x) \approx \int_1^{x+1} i^{-\beta} di,$$

and that $s(x)$ is the left-endpoint Riemann sum. Notice that f is a strictly decreasing function for positive i . As a result, we know that the left-endpoint Riemann sum must be greater than or equal to the value of the appropriate integral, which is to say that

$$s(x) \geq \int_1^{x+1} i^{-\beta} di = \frac{1}{1-\beta} ((x+1)^{1-\beta} - 1).$$

Thus, tying this all together—particularly by utilizing the fact that the exponential is an increasing function and thus maintains inequalities—shows that

$$\log((1 + 1^{-\beta})(1 + 2^{-\beta}) \cdots (1 + x^{-\beta})) \geq \exp(s(x)) \geq c \exp\left(\frac{(x+1)^{1-\beta}}{1-\beta}\right),$$

where

$$c = \exp\left(-\frac{1}{1-\beta}\right).$$

Therefore, we can now estimate the value B by noting that

$$B \leq \sum_{x=1}^{\infty} \frac{2 + x^{-\beta}}{c \exp\left(\frac{(x+1)^{1-\beta}}{1-\beta}\right)} \leq \sum_{x=1}^{\infty} \frac{3}{c \exp\left(\frac{(x+1)^{1-\beta}}{1-\beta}\right)}.$$

Finally, since $0 < \beta < 1$, then the function

$$g(x) = \frac{(1+x)^{1-\beta}}{1-\beta}$$

is increasing, so that the denominator in the final sum is an exponential of an increasing function. As a result, it can be concluded that

$$\sum_{x=1}^{\infty} \frac{3}{c \exp\left(\frac{(x+1)^{1-\beta}}{1-\beta}\right)} < \infty,$$

which shows that B is finitely-valued. Therefore, when $\beta < 1$, the birth and death chain is positive-recurrent.

1.2 Chains with Variable Step Size, and an Investigation of Hitting Times

Unfortunately, recurrence can only paint so much of the picture of a system's fragility. Particularly, birth and death chains can be used to model real world scenarios, and the parameters governing these scenarios are often subject to perturbation. As a result, an important part of the conception of mathematical fragility being developed here is understanding how the birth and death chain reacts to perturbations in its parameters. The most natural parameter to investigate first is the size of the step when transitioning away from a particular state. Further, it is not true generally that all states of a birth and death chain share the same classification, and so a more general measurement of the chain's behavior is required. These measurements are *hitting times*.

Definition 1.2.1 Suppose that X_n is a Markov chain with state space \mathcal{S} . The *hitting time* from $i \in \mathcal{S}$ to $j \in \mathcal{S}$ is the expected amount of time required for the chain to achieve state j given that its initial position is i .

In order to investigate how hitting times respond to perturbations in step size, it is first helpful to simplify the scenario, imposing that the state space be \mathbb{N} and that each step has size s . Further, let $h_s(x)$ be the hitting time to 0 given an initial position $X_0 = x$, under the fixed step size s . Returning to the metaphor of a walker, the quantity $h_s(x)$ should represent the expected amount of time for a person initially positioned at time x to reach the origin assuming a step size of s . Put differently, $h_s(x)$ somehow quantifies how efficiently the chain can move from x to 0.

As such, there are a number of quantities of interest involved in the investigation of this particular system's fragility:

- $h_s(x)$ — the expected amount of time for the chain to hit 0 from initial position x under a fixed step size s .
- $\frac{\partial h_s(x)}{\partial s}$ — the instantaneous change in hitting time to 0 with respect to change in the step size s . If this is a negative quantity, it is expected that increases in step size should increase the efficiency of the chain reaching 0 from its initial position. In contrast, if this is a positive quantity then less efficiency is expected.
- $\frac{h_s(x)}{x/s}$ — a ratio of the hitting time to 0 as compared to the minimum number of steps required to reach 0 from initial position x .
- $\frac{h_s(x)}{x}$ — a ratio of the hitting time to 0 as compared to a "neutral" scenario in which the expected number of steps to reach 0 is exactly the position x .

I began investigation of hitting times by first looking for a birth and death chain in which the expected hitting time from any n is given by $h_s(n) = n/s$. This is, in some sense, what is expected to be the most minimal scenario. Doing so involved the following relationship:

$$h_s(x) = p(x)(1 + h_s(x + s)) + q(x)(1 + h_s(x - s)).$$

This can be thought of as saying that the expected hitting time from x is equal to the probability of stepping in the positive direction times the associated hitting time, plus the probability of stepping in the negative direction times its associated hitting time. However, we also know that the associated hitting times should be one more than the hitting times from the new positions, which shows that the associated hitting times are $1 + h_s(x + s)$ and $1 + h_s(x - s)$, respectively. In fact, this holds true for all of the random walks considered here, and will be used in further investigation.

Returning now to the scenario at hand, we can begin to deduce further information about a random walk with such a property. By substitution in the given relationship using the expected hitting times, we have that

$$\begin{aligned} \frac{x}{s} &= p(x) \left(1 + \frac{x+s}{s} \right) + q(x) \left(1 + \frac{x-s}{s} \right) \\ &= p(x) + q(x) + p(x) \frac{x+s}{s} + q(x) \frac{x-s}{s} \\ &= 1 + p(x) \frac{x+s}{s} + q(x) \frac{x-s}{s}. \end{aligned}$$

Therefore,

$$\begin{aligned} x &= s + p(x)(x + s) + q(x)(x - s) \\ &= s + p(x)(x + s) + (1 - p(x))(x - s) \\ &= s + x - s + p(x)(x + s - (x - s)) \\ &= x + 2s p(x). \end{aligned}$$

Thus, we deduce that $p(x) = 0$ should have this property. In particular, this says that $q(x) = 1$, in which case this is the expected result. The random walk that moves in the negative direction with probability 1 should take the minimal amount of time to reach 0.

This will allow us to make interesting comparisons about other random walks that we investigate. For example, it may be interesting to find a random walk where the expected return time has no dependence on the size of the step. A reasonable place to begin such investigation is by looking for a chain which produces the hitting time $h_s(x) = x$. We begin again by substitution into the starting relationship.

$$\begin{aligned}
 x &= p(x)(1+x+s) + q(x)(1+x-s) \\
 &= p(x) + q(x) + p(x)(x+s) + q(x)(x-s) \\
 &= 1 + p(x)(x+s) + (1-p(x))(x-s) \\
 &= 1 + x - s + p(x)(x+s - (x-s)) \\
 &= 1 + x - s + 2s p(x).
 \end{aligned}$$

Thus, we deduce that $s - 1 = 2s p(x)$. Therefore we find that

$$p(x) = \frac{s-1}{2s}.$$

This seemed like a reasonable place to run some new simulations investigating this family of probability distributions and their relation to step-size.